deeplearning.ai

Case Studies

Why look at case studies?

Outline

Classic networks:

- LeNet-5 ⇐
- AlexNet ⇐
- VGG ⇐

ResNet (152)

Inception ⇒

NN های کد قرار است بحث کند.

deeplearning.ai

Case Studies

Classic networks

LeNet - 5

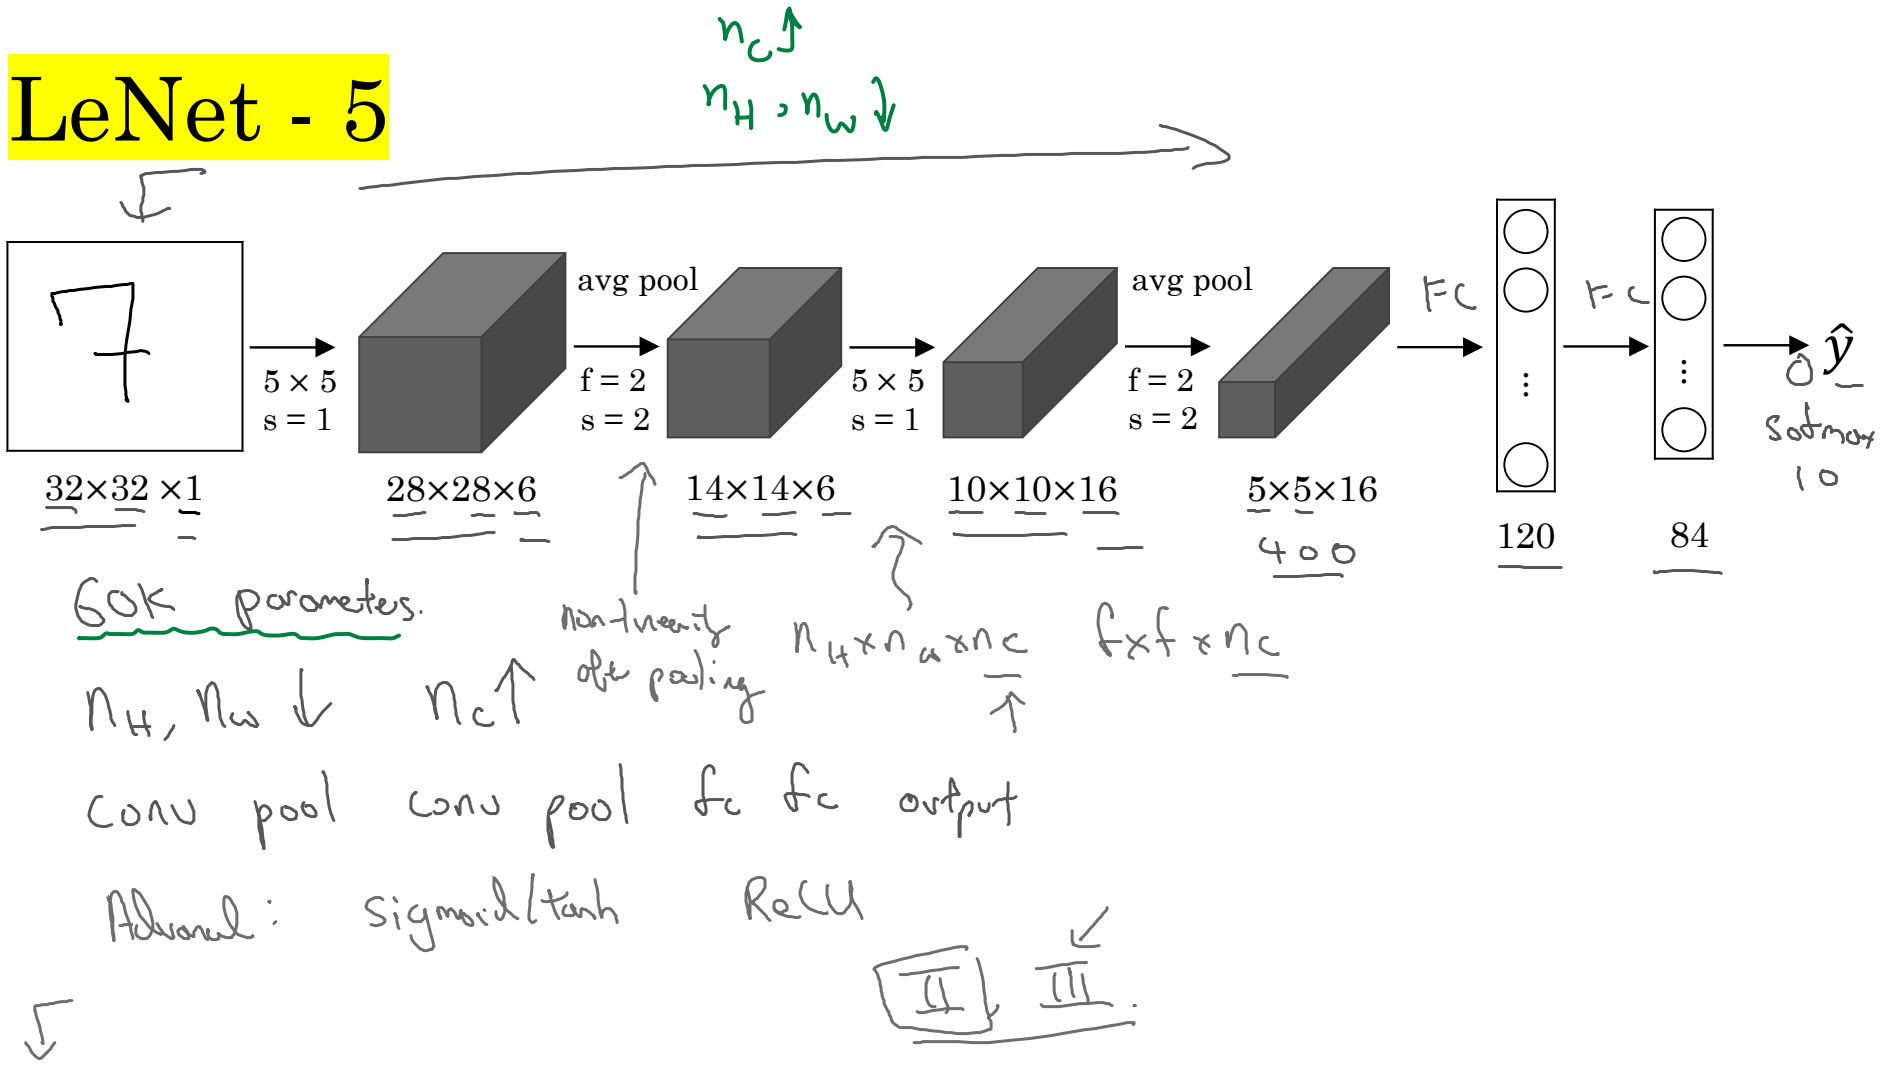

[LeCun et al., 1998. Gradient-based learning applied to document recognition]

AlexNet

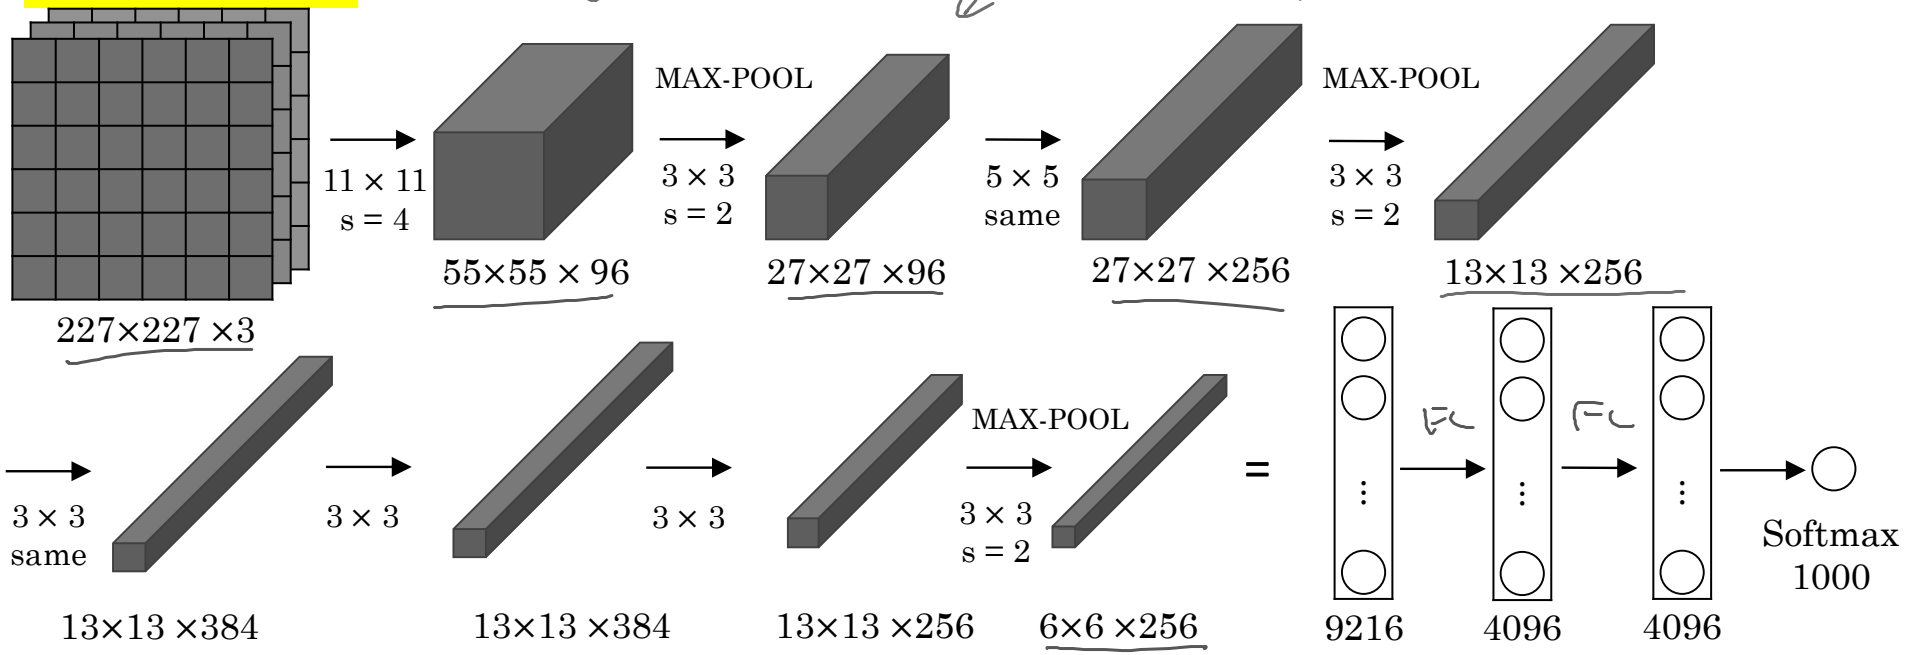

- Similar to LeNet, but much bigger.
- ReLU
- Multiple GPUs.
- Local Response Normalization (LRN)

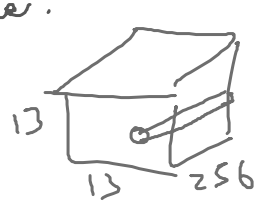

60M parameters

[Krizhevsky et al., 2012. ImageNet classification with deep convolutional neural networks]

VGG - 16

CONV = 3×3 filter, s = 1, same

MAX-POOL = 2×2 , s = 2

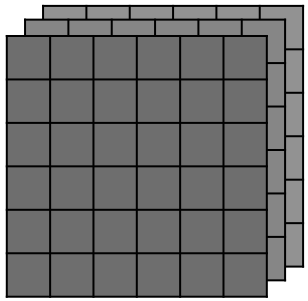

224×224 ×3

[CONV 64] ×2

224×224×64

POOL

112×112 ×64

[CONV 128] ×2

112×112 ×128

POOL

56×56 ×128

[CONV 256] ×3

56×56 ×256

POOL

28×28 ×256

[CONV 512] × 3

28×28 ×512

POOL

14×14×512

[CONV 512] ×3

14×14 ×512

POOL

7×7×512

FC 4096

FC 4096

Softmax 1000

[Simonyan & Zisserman 2015. Very deep convolutional networks for large-scale image recognition]

deeplearning.ai

Case Studies

Residual Networks (ResNets)

Residual block

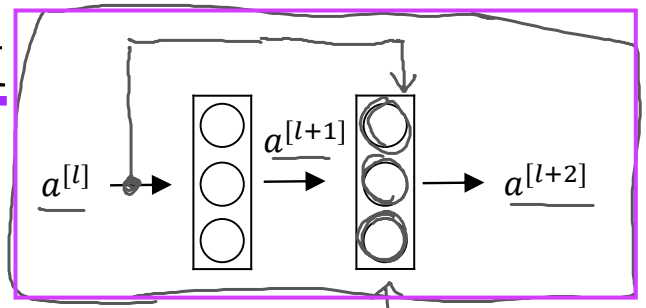

"short cut" / skip connection

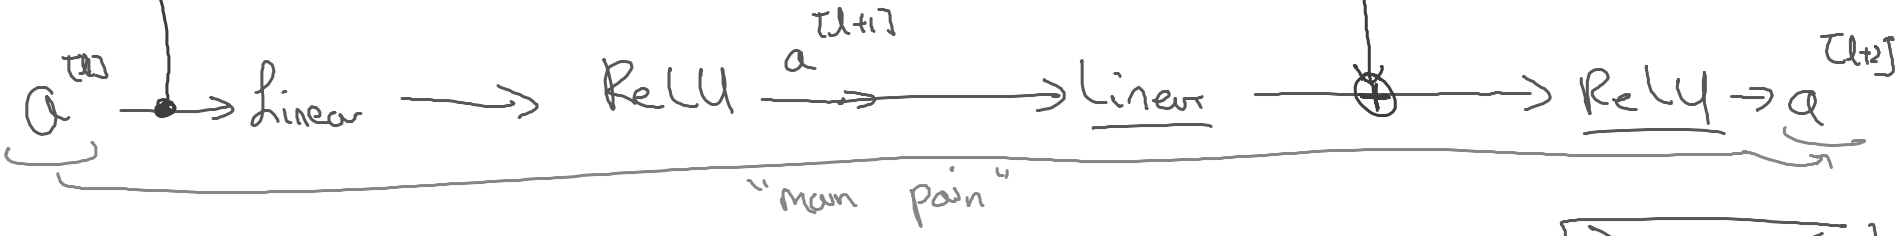

$$z^{[l+1]} = W^{[l+1]} a^{[l]} + b^{[l+1]} \qquad a^{[l+1]} = g(z^{[l+1]}) \qquad z^{[l+2]} = W^{[l+2]} a^{[l+1]} + b^{[l+2]} \qquad \cancel{a^{[l+2]} = g(z^{[l+2]})}$$

$$a^{[l+2]} = g(z^{[l+2]} + a^{[l]})$$

[He et al., 2015. Deep residual networks for image recognition]

Residual Network

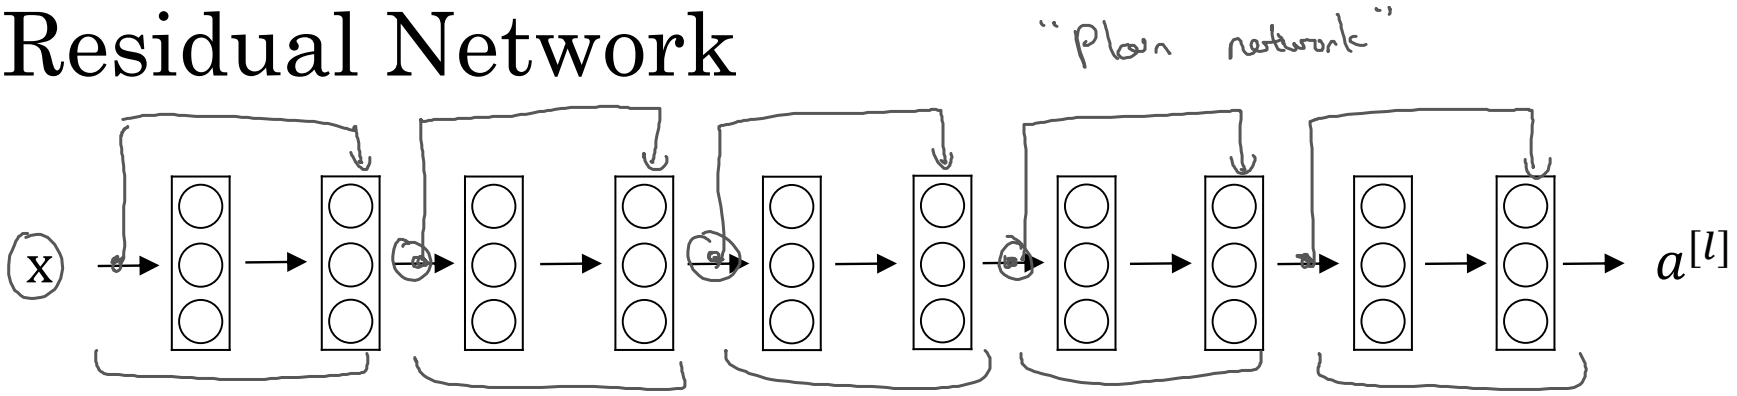

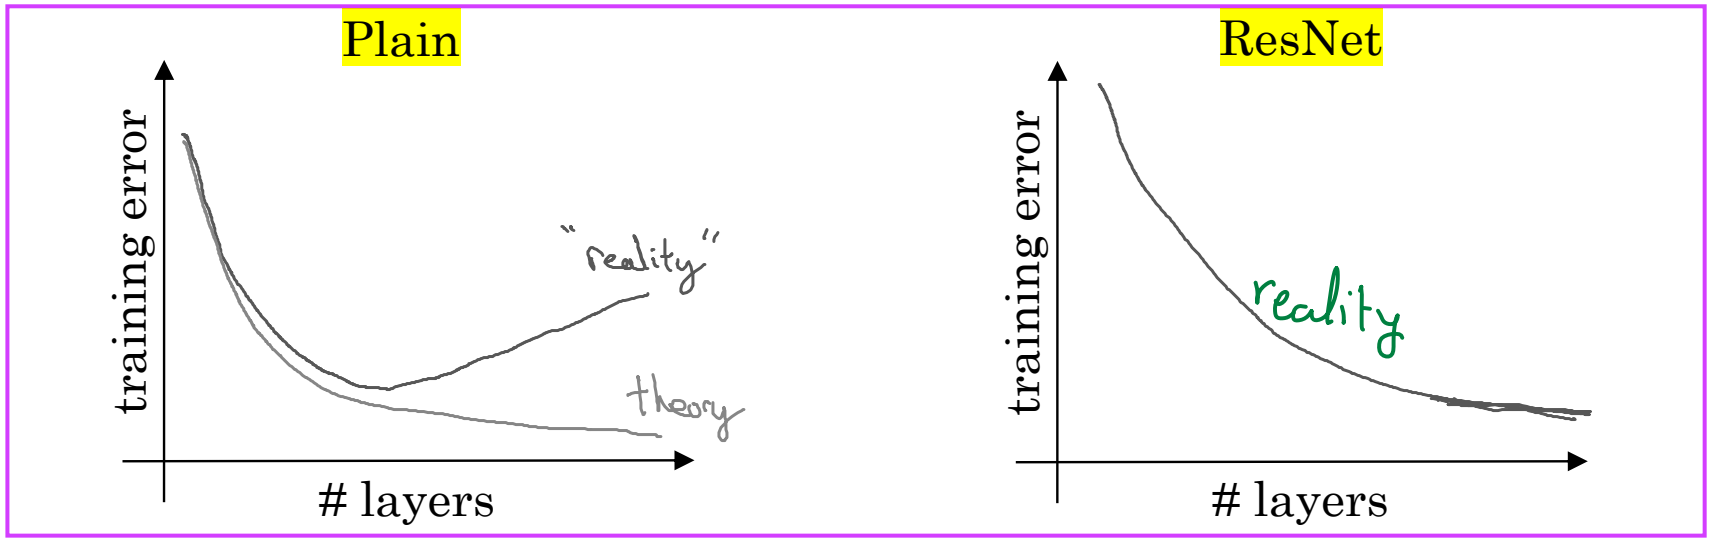

[He et al., 2015. Deep residual networks for image recognition]

deeplearning.ai

Case Studies

Network in Network and 1×1 convolutions

$28 \times 28 \times 192 \xrightarrow[1 \times 1 \times 192]{\text{Conv } 1 \times 1} 28 \times 28 \times n_c$ (# filters)

Why does a 1 × 1 convolution do?

changes n_c, fixes n_H, n_W

1	2	3	6	5	8
3	5	5	1	3	4
2	1	3	4	9	3
4	7	8	5	7	9
1	5	3	7	4	8
5	4	9	8	3	5

$6 \times 6 \times 1$

* [2] = | 2 | 4 | 6 | ... | | |

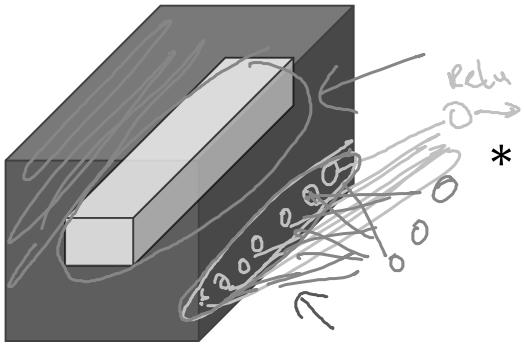

$6 \times 6 \times 32$

*

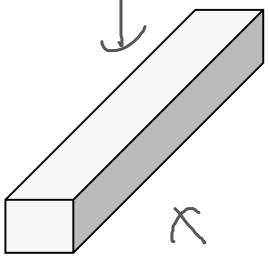

$1 \times 1 \times 32$

32 → # filters. $n_c^{[l+1]}$

=

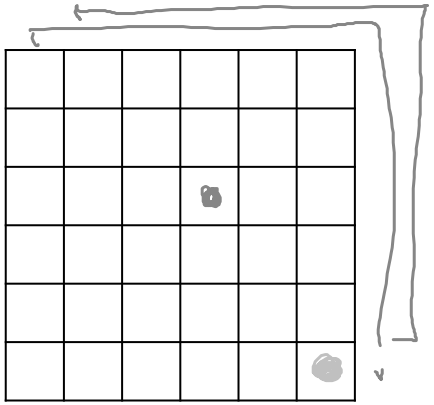

$6 \times 6 \times$ # filters

ReLU

Network in Network

[Lin et al., 2013. Network in network]

حتی می‌توان تبدیل 28×28×192 → 28×28×192 داشت که در این حالت نقش conv ۱×۱ اضافه کردن non-linearity است.

Using 1×1 convolutions

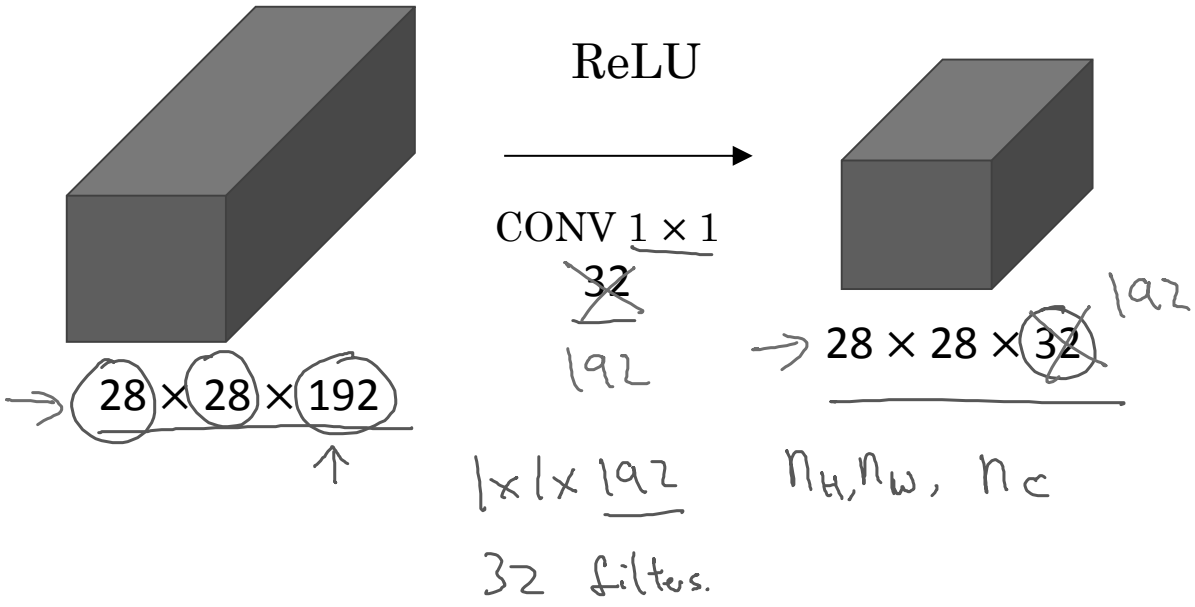

28 × 28 × 192 → 28 × 28 × 32 برای آن به 32 تا conv ۱×۱ که هرکدام ۱×۱×192 هستند احتیاج دارم.

[Lin et al., 2013. Network in network]

deeplearning.ai

Case Studies

Inception network motivation

28×28×192 → 28×28×256

Motivation for inception network

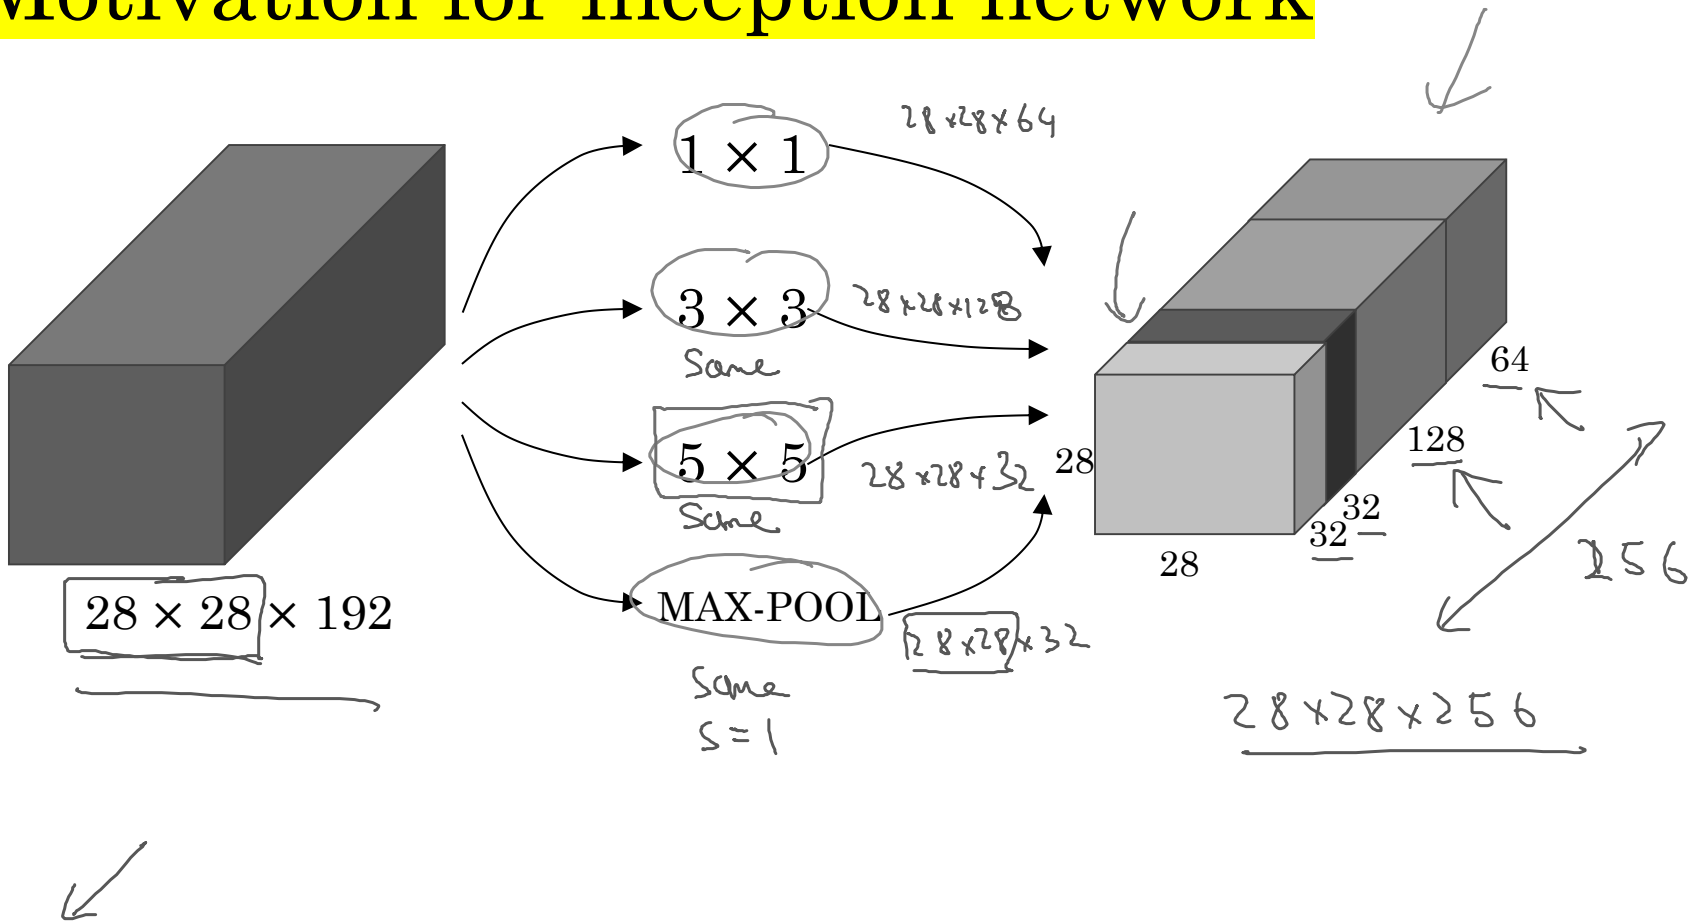

[Szegedy et al. 2014. Going deeper with convolutions]

The problem of computational cost

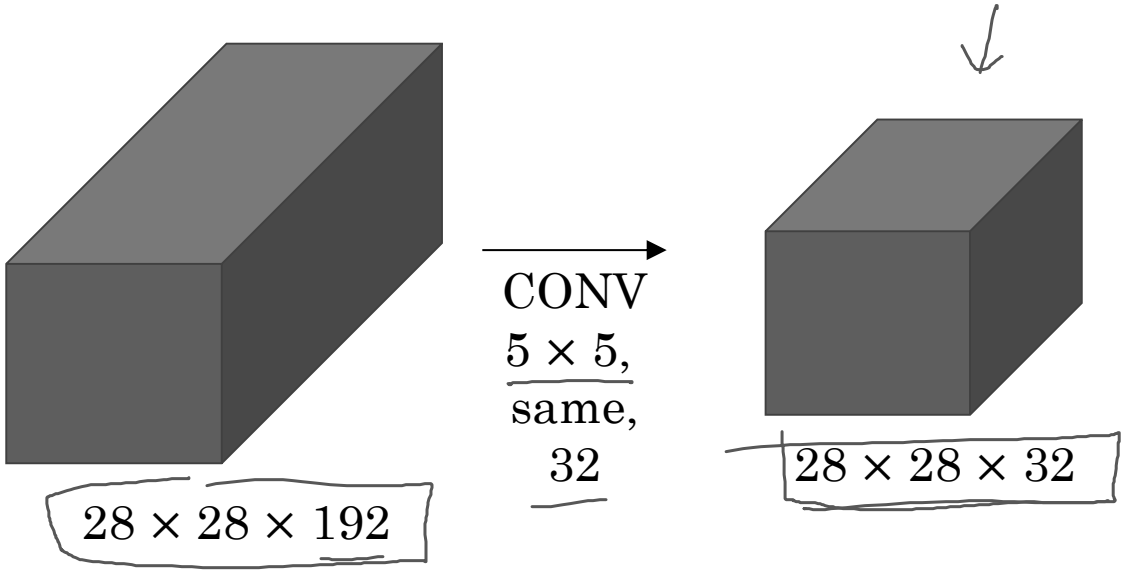

32 filters. filter one 5x5x192.

28x28x32 x 5x5x192 = 120M.

Using 1×1 convolution

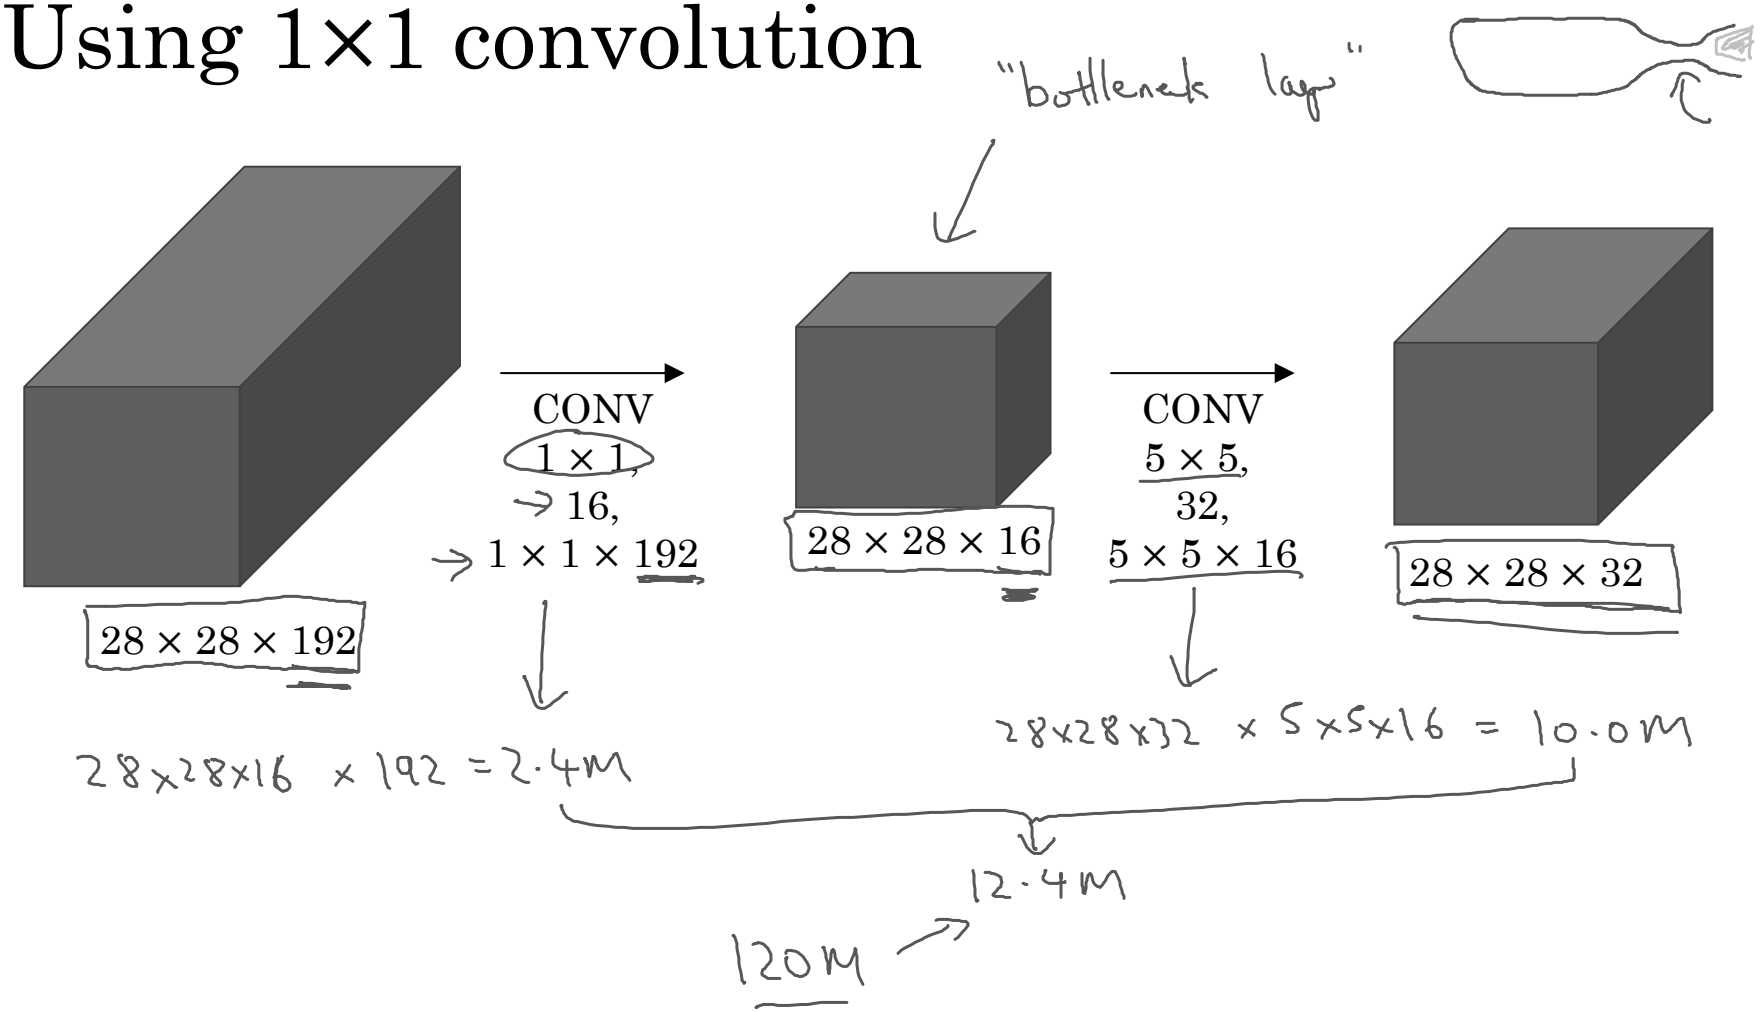

deeplearning.ai

Case Studies

Inception network

Inception module

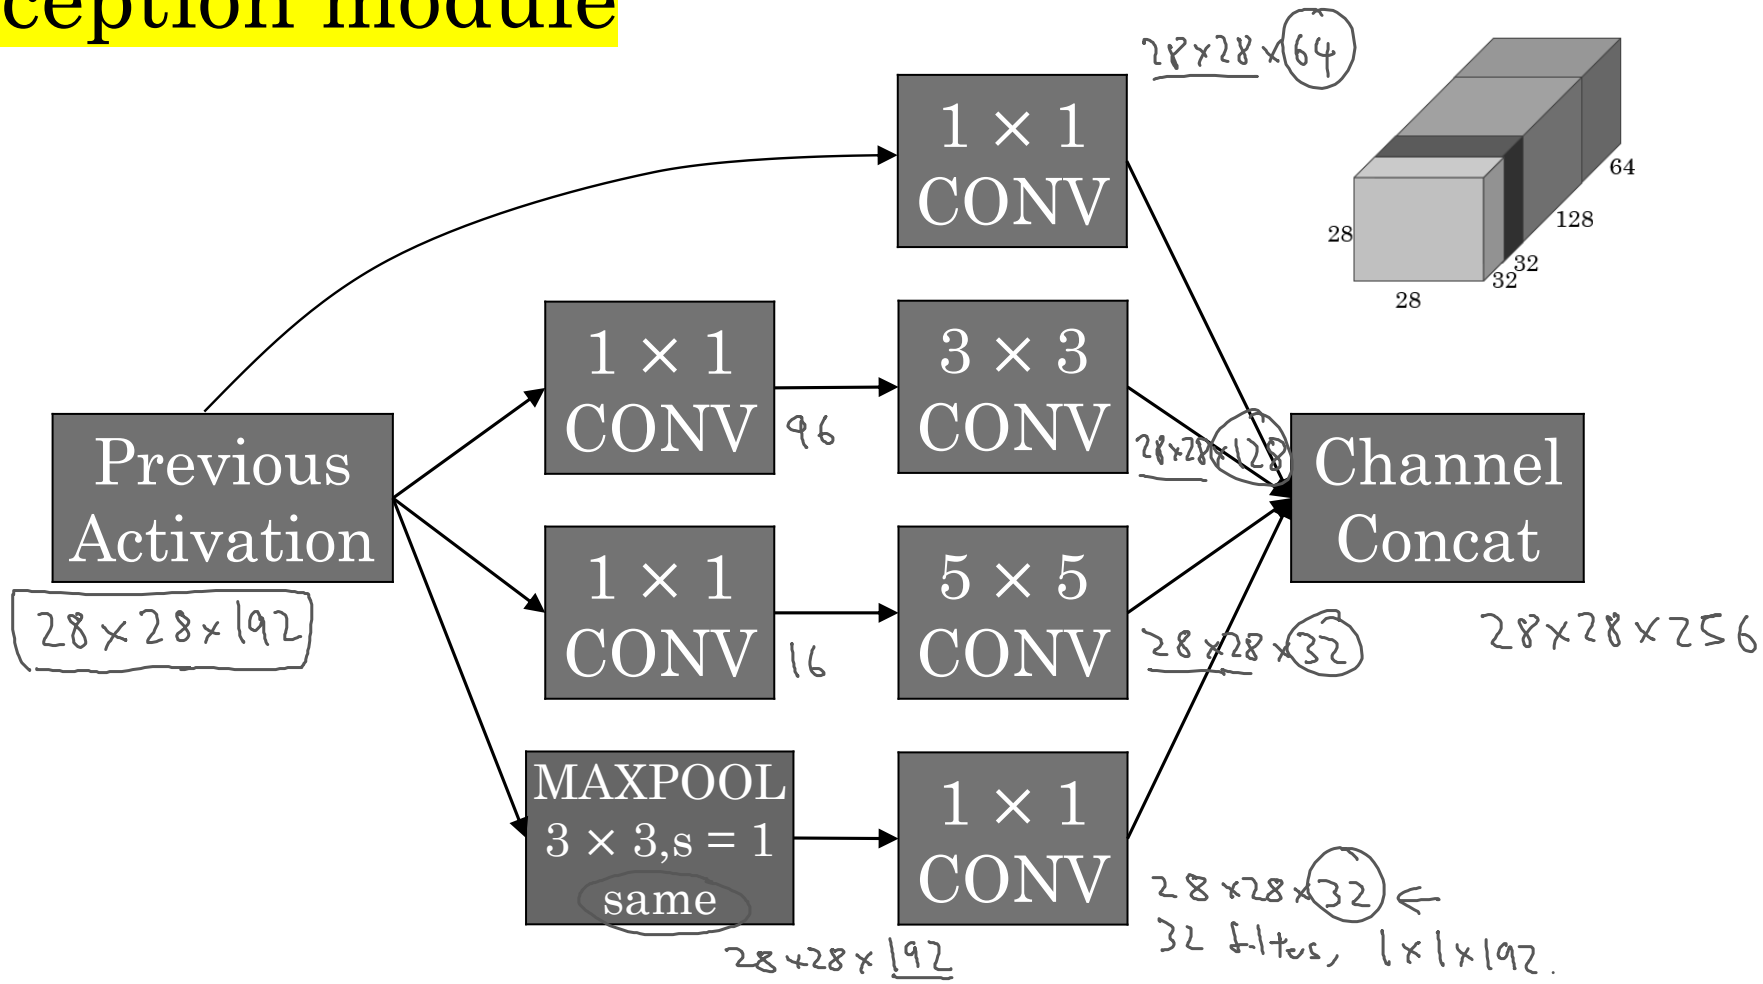

Inception network

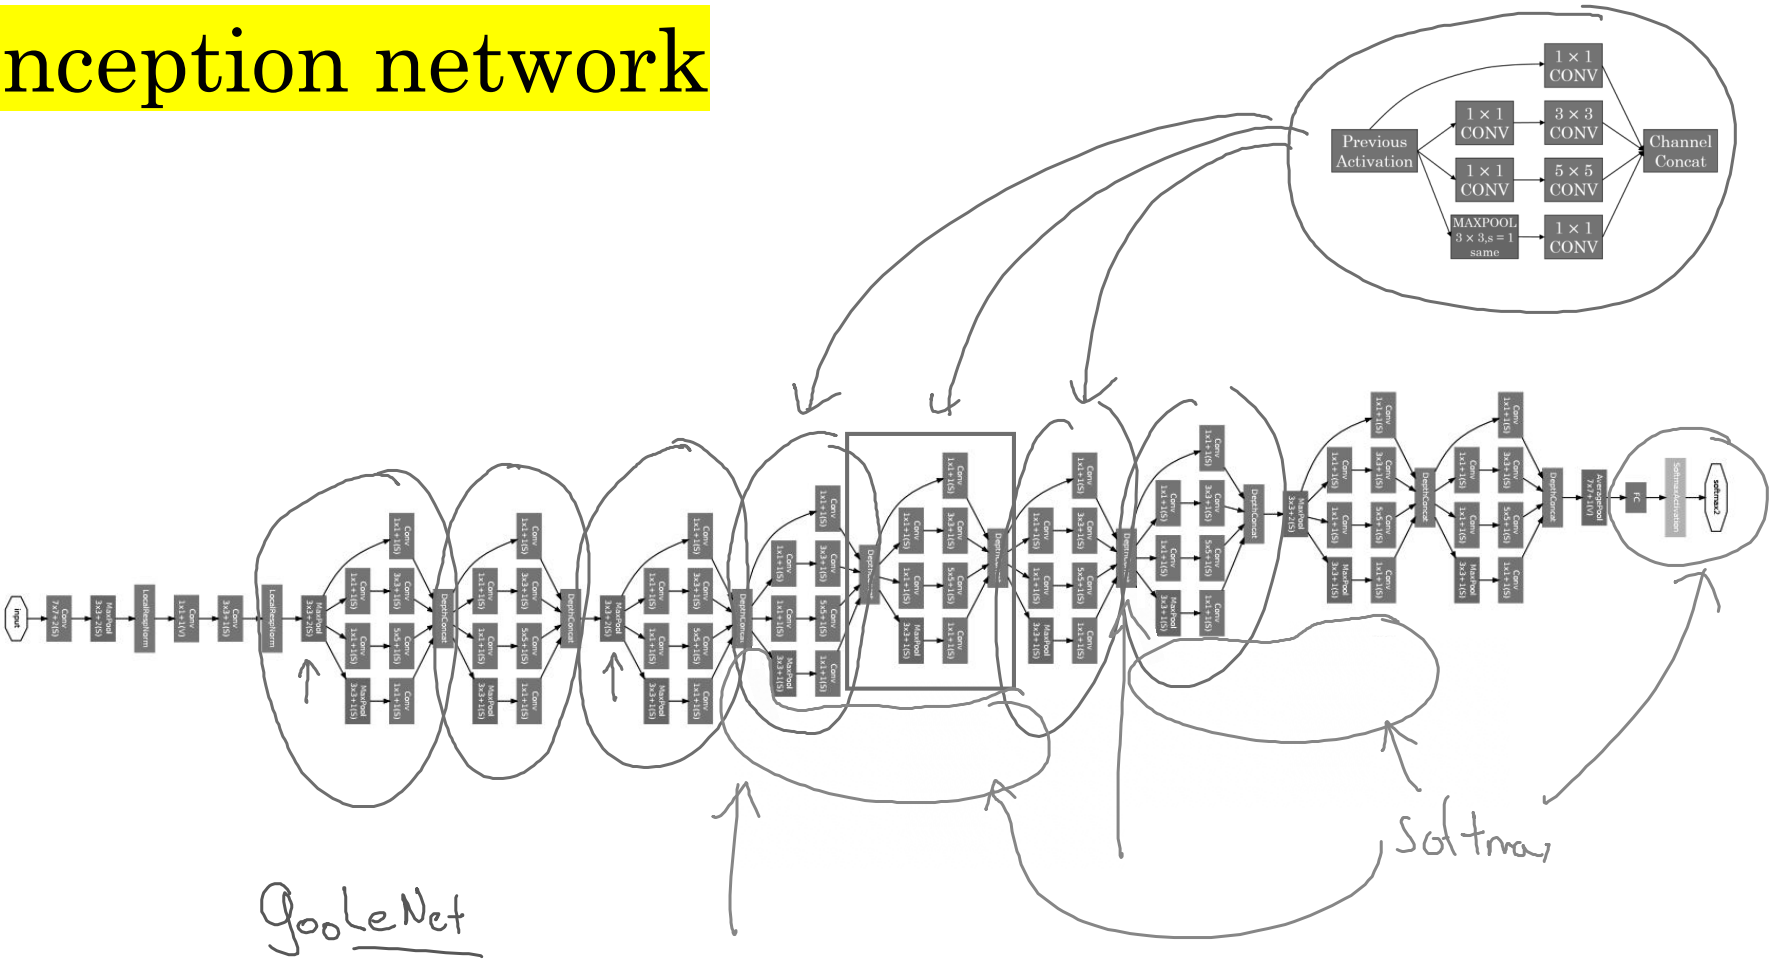

[Szegedy et al., 2014, Going Deeper with Convolutions]

http://knowyourmeme.com/memes/we-need-to-go-deeper

Tips for doing well on benchmarks/winning competitions

Ensembling

3-15 networks

→ $\hat{y}$

- Train several networks independently and average their outputs

Multi-crop at test time

- Run classifier on multiple versions of test images and average results

10-crop

Use open source code

- Use architectures of networks published in the literature
- Use open source implementations if possible
- Use pretrained models and fine-tune on your dataset